

Side Channels and Deep Neural Network Weights

Attacks, Defences and the Future to Come

—
Journées GT SSLR 2025

Lorenzo CASALINO

CentraleSupélec, IRISA, Inria (SUSHI Team)

lorenzo.casalino@{centralesupelec.fr, irisa.fr, inria.fr}



ANR-21-CE39-0018

Agenda

1 Background and Motivation

2 Attacks: Methodologies and Challenges

- Single Neuron
- Whole Network

3 Defences: Methodologies and Challenges

- Masking
- Shuffling
- Other Approaches

4 Conclusions

Table of Contents

1 Background and Motivation

2 Attacks: Methodologies and Challenges

- Single Neuron
- Whole Network

3 Defences: Methodologies and Challenges

- Masking
- Shuffling
- Other Approaches

4 Conclusions

Side-channel Analysis




Figure: Information Recovery Through Side-channel Analysis

Deep Neural Networks (DNNs)




Figure: A simple DNN brand classifier¹.




Figure: A neuron computes a weighted sum of its inputs (Eq. 1).

¹Duck and Kurzgesagt Logo belong to Kurzgesagt

Motivation




Figure: The training process.²

DNN Training is Expensive

- Expensive hardware (e.g., GPUs), time-intensive (e.g., days)

Weights Piracy

A non-negligible **economic damage**

²Duck (Kurzgesagt), Shrimp (Jellycat London), Masterball (Nintendo)

Table of Contents

- 1 Background and Motivation
- 2 Attacks: Methodologies and Challenges
 - Single Neuron
 - Whole Network
- 3 Defences: Methodologies and Challenges
 - Masking
 - Shuffling
 - Other Approaches
- 4 Conclusions

Recovery of Weights – Single Neuron

Weight Recovery Attack

Retrieve **correct** weight value *among all* the accepted ones.

$$o_6 = \underbrace{i_0 \cdot w_0}_{\text{target}} + i_1 \cdot w_1 + i_2 \cdot w_2$$

target oper.

Attack Complexity for a Neuron

- **Attack Complexity:** $O(N_{weights})$, $N_{weights} = \#\text{weights}$
 - Typical $N_{weights}$: 9 (MobileNet-v2), 25 (GoogleLeNet)
 - Already a non-negligible effort
- But actually ...

Recovery of Weights – Single Neuron

Weights and Data Types

- Weights data type: floating-point or integer
- Weights may have wider or narrower bitwidths (e.g., 32 bits)
- For each data type and bitwidth, attack strategies and complexities change

Work	Type/Width	Complexity (at least)
[Jou+23]	Float/32	$O(2^{16} \cdot N_{\text{weights}})$
[Yos+21]	Int/8	$O(2^{8k} + N_{\text{weights}})$
[Gon+24]	Int/8	$O(2^{16} + N_{\text{weights}})$

Table: Complexity of State-of-the-Art Weight Recovery Attacks (One Neuron).

Recovery of Weights – Whole Network

Attacking the Whole Network

- Attacker can independently target neurons (of the same layer)
- Attack cost linear with number of neurons (N_{neurons})
- DNNs with millions of neurons \implies millions weights ($N_{\text{weights,net}}$)
 - Examples: $\sim 3,4M$ (MobileNet-v2), $\sim 6,8M$ (GoogleLeNet)

Work	Type/Width	Complexity
[Jou+23]	Float/32	$O(2^{16} \cdot N_{\text{weights,net}})$
[Yos+21]	Int/8	$O((2^{8k} \cdot N_{\text{neurons}} + N_{\text{weights,net}}))$
[Gon+24]	Int/8	$O((2^{16} \cdot N_{\text{neurons}} + N_{\text{weights,net}}))$

Table: Complexity of State-of-the-Art Weight Recovery Attacks (Whole Network).

Attacks: to Sum Up

Weight Recovery – A Challenging Task

- Weight recovery linear in number of weights

Example:

- DNN with 600k weights (in total)³
- Weight Recovery Time: 10 seconds/weight
- **Recovery time: 69 days**
- Hidden constants increase the recovery time
- Not considering other costs (e.g., side-channel acquisition time)
- Attacking beyond input layer adds further difficulty

State-of-the-Art Limitations

- Methodologies proved only on really small networks
- Very few works target beyond input layer


³Reasonable for microcontroller-oriented DNNs
(<https://github.com/mit-han-lab/mcunet>)

Table of Contents

- 1 Background and Motivation
- 2 Attacks: Methodologies and Challenges
 - Single Neuron
 - Whole Network
- 3 Defences: Methodologies and Challenges
 - Masking
 - Shuffling
 - Other Approaches
- 4 Conclusions

Masking

Masking



Masking

Replace the weight-dependent signal with N random ones (the *shares*)

Advantages

Provably secure side-channel countermeasures

Difficulties

- ① Slower, huge (code size/silicon area), and energy-ravenous design
- ② Physical non-idealities may lead to information leakage [Cas+23]
- ③ Huge design limits security evaluation

Shuffling

Inference #0 : $i_0 \cdot w_0 + i_1 \cdot w_1 + i_2 \cdot w_2$

Inference #1 : $i_1 \cdot w_1 + i_2 \cdot w_2 + i_0 \cdot w_0$

Inference #2 : $i_2 \cdot w_2 + i_0 \cdot w_0 + i_1 \cdot w_1$

Shuffling

- Randomly shuffle operations to bury weight-dependent signal in signal noise

Advantages

Less expensive than masking

Difficulties

- 1 No formal security guarantees
- 2 Operations (e.g., division) may lead to unintended information leakage [Puš+25]
- 3 No generic security projections (attacker dependent)



Other Approaches

DNN-Tailored Countermeasures

- Current defences come from cryptanalysis
- But DNNs \neq cryptosystems!
- DNNs exhibit particular characteristics (e.g., error resilience)

Approximate-Computing (AxC)-based Countermeasures

- Trade accuracy for better energy efficiency, size and execution time
- Recently considered as a countermeasure [Din+25; Jap+25][Cas+26]⁴

⁴Paper just accepted at HOST'26

Table of Contents

- 1 Background and Motivation
- 2 Attacks: Methodologies and Challenges
 - Single Neuron
 - Whole Network
- 3 Defences: Methodologies and Challenges
 - Masking
 - Shuffling
 - Other Approaches
- 4 Conclusions

Challenges

Attack Methodologies

- Linear complexity with #weights
 - But million of weights
 - Non-negligible hidden constants
- No attempts on full DNN models

Defence Methodologies

- Too expensive to deploy, design and evaluate (masking)
- Provide few security guarantees (shuffling)
- Few works proposing countermeasures
- Few security analyses of countermeasures

Narrow Set of Targets

- Most works consider really simple MLPs and CNNs
- No attempts on state-of-the-art DNN models
- Marginal focus on other NNs (e.g., Spiking NNs [PBS25]).

The Future to Come

Better Evaluation Methodologies

- Efficient and Comprehensive (e.g., analyse deeper layers, use all leaked information)
- Explainable (i.e., precisely identify the leakage root cause)

↓ to have ↑

Better Defence Methodologies

- Efficient (i.e., minimise performance overhead)
- Effective (i.e., protect against state-of-the-art attacks)

That's All Folks

Thank You!

Bibliography I

- [Bro+24] Manuel Brosch et al. "A Masked Hardware Accelerator for Feed-Forward Neural Networks With Fixed-Point Arithmetic". In: *IEEE VLSI* (2024).
- [Cas+23] Lorenzo Casalino et al. "A Tale of Resilience: On the Practical Security of Masked Software Implementations". In: *IEEE Access* (2023).
- [Cas+26] Lorenzo Casalino et al. "Double-Strike: Breaking Approximation-based Side-Channel Countermeasures for DNNs". In: *HOST* (2026).
- [Din+25] Ruyi Ding et al. "MACPruning: Dynamic Operation Pruning to Mitigate Side-Channel DNN Model Extraction". In: *HOST* (2025).
- [Gon+24] Cheng Gongye et al. "Side-Channel-Assisted Reverse-Engineering of Encrypted DNN Hardware Accelerator IP and Attack Surface Exploration". In: *IEEE S&P.* 2024.

Bibliography II

- [Jap+25] Aditya Japa et al. “Security of Approximate Neural Networks against Power Side-channel Attacks”. In: *IEEE DAC*. 2025.
- [Jou+23] Raphaël Joud et al. “A Practical Introduction to Side-Channel Extraction of Deep Neural Network Parameters”. In: *Smart Card Research and Advanced Applications*. 2023.
- [PBS25] Matthias Probst, Manuel Brosch, and Georg Sigl. “Side-Channel Analysis of Integrate-and-Fire Neurons Within Spiking Neural Networks”. In: *IEEE Transactions on Circuits and Systems I* (2025).
- [Puš+25] Leonard Puškáč et al. “Make Shuffling Great Again: A Side-Channel-Resistant Fisher–Yates Algorithm for Protecting Neural Networks”. In: *IEEE VLSI* (2025).

Bibliography III

- [Yos+21] Kota Yoshida et al. “Model Reverse-Engineering Attack against Systolic-Array-Based DNN Accelerator Using Correlation Power Analysis”. In: *IEICE Transactions on Fundamentals of Electronics, Communications and Computer Sciences* (2021).

Recovery of Weights – Hidden Layers




Figure: σ Influences Next Layer's Inputs.

$$\sigma(x) = \begin{cases} x & x \geq 1 \\ 0 & \text{otherwise} \end{cases}$$

Figure: ReLU
Activation Function


Attacker Needs Full Input Control

- Hidden layer's input depends on previous layer
- This dependency may forbid hidden layers' weight recovery

Example:

- $o_{1,4} = -1.4 \rightarrow \sigma(o_{1,4}) = 0 \rightarrow \sigma(o_{1,3}) \cdot w_{0,4} = 0$
- Cannot attack $w_{1,4}$!

Recovery of Weights – Hidden Layers



$$\sigma(x) = \begin{cases} x & x \geq 1 \\ 0 & \text{otherwise} \end{cases}$$

Figure: ReLU Activation Function

Figure: σ influence Next Layer's Inputs.

State-of-the-art Solutions [Gon+24; PBS25]

- **Idea:** determine inputs i_j to control $\sigma(o_{h,k})$ (hidden layer's inputs)
- **Caveat:** attack complexity \propto Inputs per layer, Number of layers

Masking – More on the Cost

Table: Software Masked CNN – Execution Time Overheads (Excerpt from [Bro+24]).

Architecture	Masked	Masked (Improved)
(6,5)-(16,5)-256-120-84	× 703%	× 238%
(16,5)-(32,5)-1568	× 306%	× 135%

Table: Software Masked CNN – Minimal Storage Requirement (Architectures from [Bro+24]).

Architecture	Original (KBytes)	Masked (2 shares, KBytes)
(6,5)-(16,5)-256-120-84	3,940	7,880
(16,5)-(32,5)-1568	11,072	22,144

Minimal Storage Requirements

$$N_{\text{weights, net}} \cdot N_{\text{shares}}$$